

K8s Service-Mesh Introduction

Wen Zhenglin
2018-11-28

概览

- 什么是Service-mesh
- Service-mesh的主要解决方案
- Istio的介绍及示例
- Linkerd简要介绍
- 总结

什么是service-mesh

Service mesh is a network of microservices that make up such applications and the interactions between them.

As a service mesh grows in size and complexity, it can become harder to understand and manage

Its requirements can include discovery, load balancing, failure recovery, metrics, and monitoring.

A service mesh also often has more complex operational requirements, like A/B testing, canary releases, rate limiting, access control, and end-to-end authentication.

Service-mesh解决方案

1. Istio
2. Linkerd



Istio

Istio介绍

Istio is a sidecar container implementation of the features and functions needed when creating and managing microservices. Monitoring, tracing, circuit breakers, routing, load balancing, fault injection, retries, timeouts, mirroring, access control, rate limiting, and more, are all a part of this.

Istio let's you get these benefits with no changes to your source code.

Istio的特性

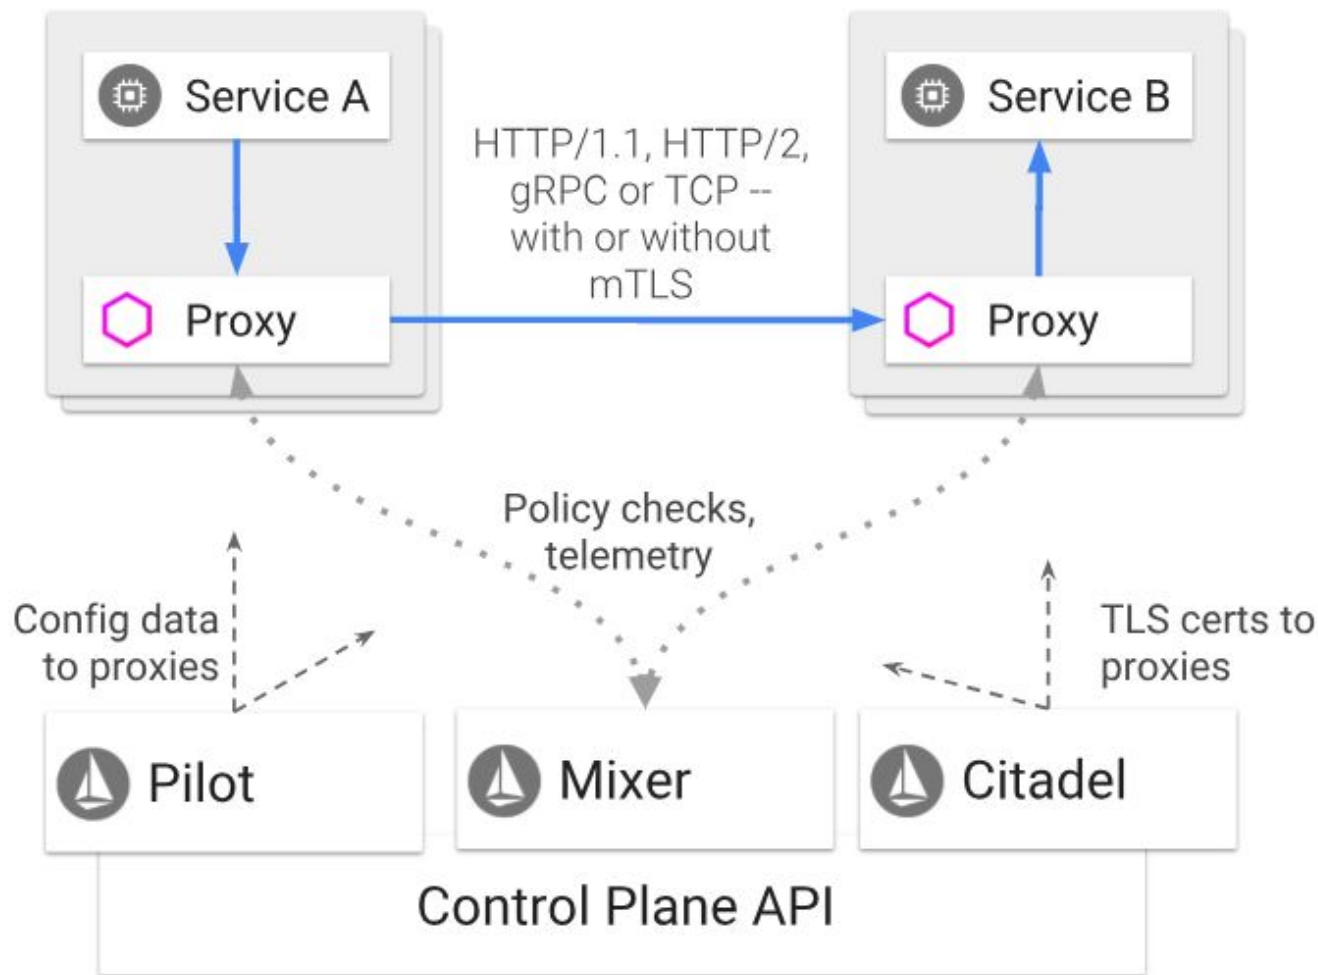
- Automatic load balancing for HTTP, gRPC, WebSocket, and TCP traffic.
- Fine-grained control of traffic behavior with rich routing rules, retries, failovers, and fault injection.
- A pluggable policy layer and configuration API supporting access controls, rate limits and quotas.
- Automatic metrics, logs, and traces for all traffic within a cluster, including cluster ingress and egress.
- Secure service-to-service communication in a cluster with strong identity-based authentication and authorization.

Istio arch

An Istio service mesh is logically split into a data plane and a control plane.

- The data plane is composed of a set of intelligent proxies ([Envoy](#)) deployed as sidecars. These proxies mediate and control all network communication between microservices along with [Mixer](#), a general-purpose policy and telemetry hub.
- The control plane manages and configures the proxies to route traffic. Additionally, the control plane configures Mixers to enforce policies and collect telemetry.

Istio arch



Envoy(proxy)提供的特性

- Dynamic service discovery
- Load balancing
- TLS termination
- HTTP/2 and gRPC proxies
- Circuit breakers
- Health checks
- Staged rollouts with %-based traffic split
- Fault injection
- Rich metrics

Rule configuration

新的四种K8s配置资源：

- A [VirtualService](#) defines the rules that control how requests for a service are routed within an Istio service mesh.
- A [DestinationRule](#) configures the set of policies to be applied to a request after VirtualService routing has occurred.
- A [ServiceEntry](#) is commonly used to enable requests to services outside of an Istio service mesh.
- A [Gateway](#) configures a load balancer for HTTP/TCP traffic, most commonly operating at the edge of the mesh to enable ingress traffic for an application.

Virtual service

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: reviews

spec:

 hosts:

 - reviews

 http:

 - route:

 - destination:

 host: reviews

 subset: v1

Destination Rule

apiVersion: networking.istio.io/v1alpha3

kind: DestinationRule

metadata:

name: reviews

spec:

host: reviews

trafficPolicy:

loadBalancer:

simple: RANDOM

subsets:

- name: v1

labels:

version: v1

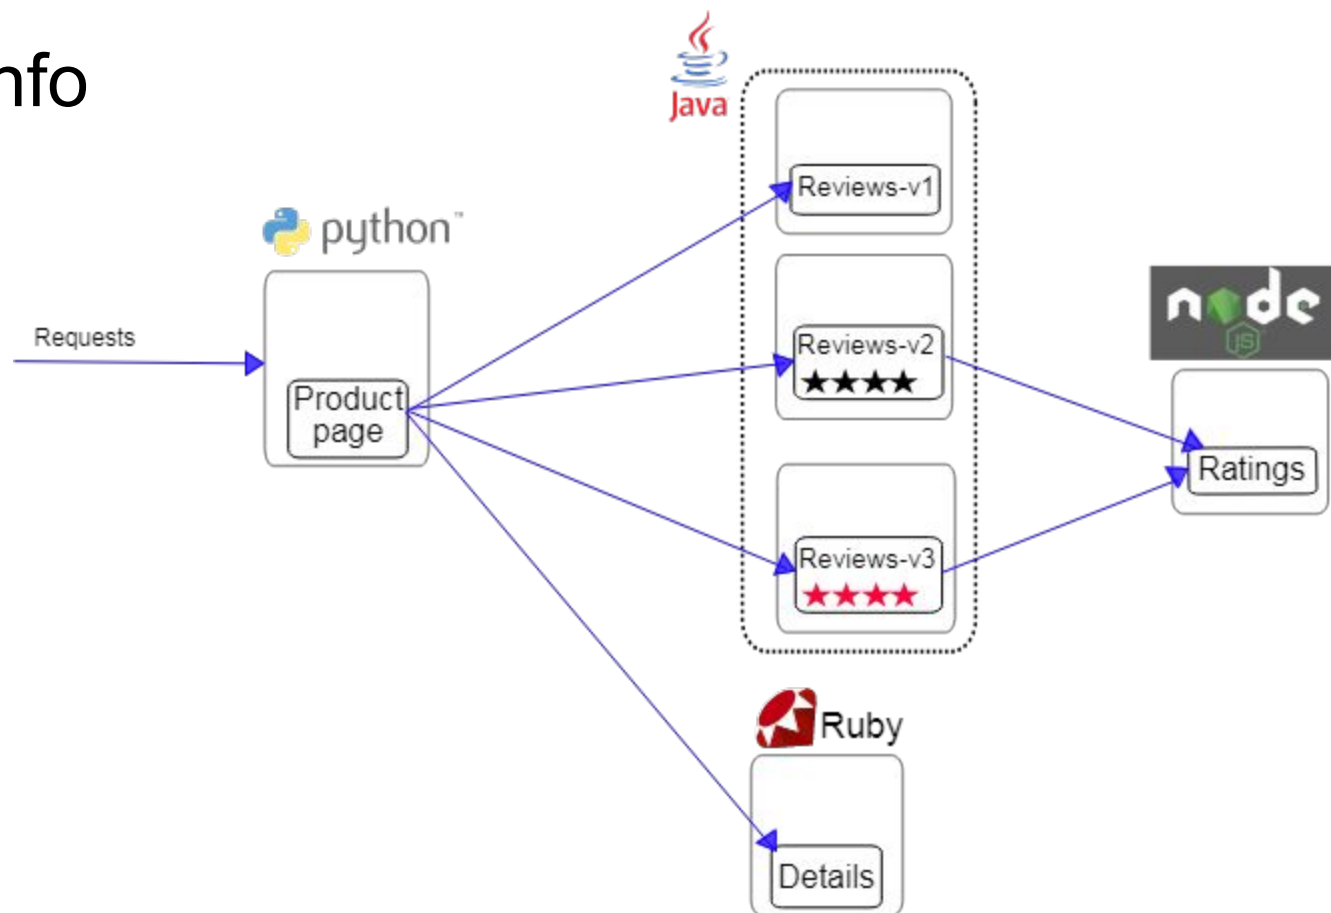
- name: v2

labels:

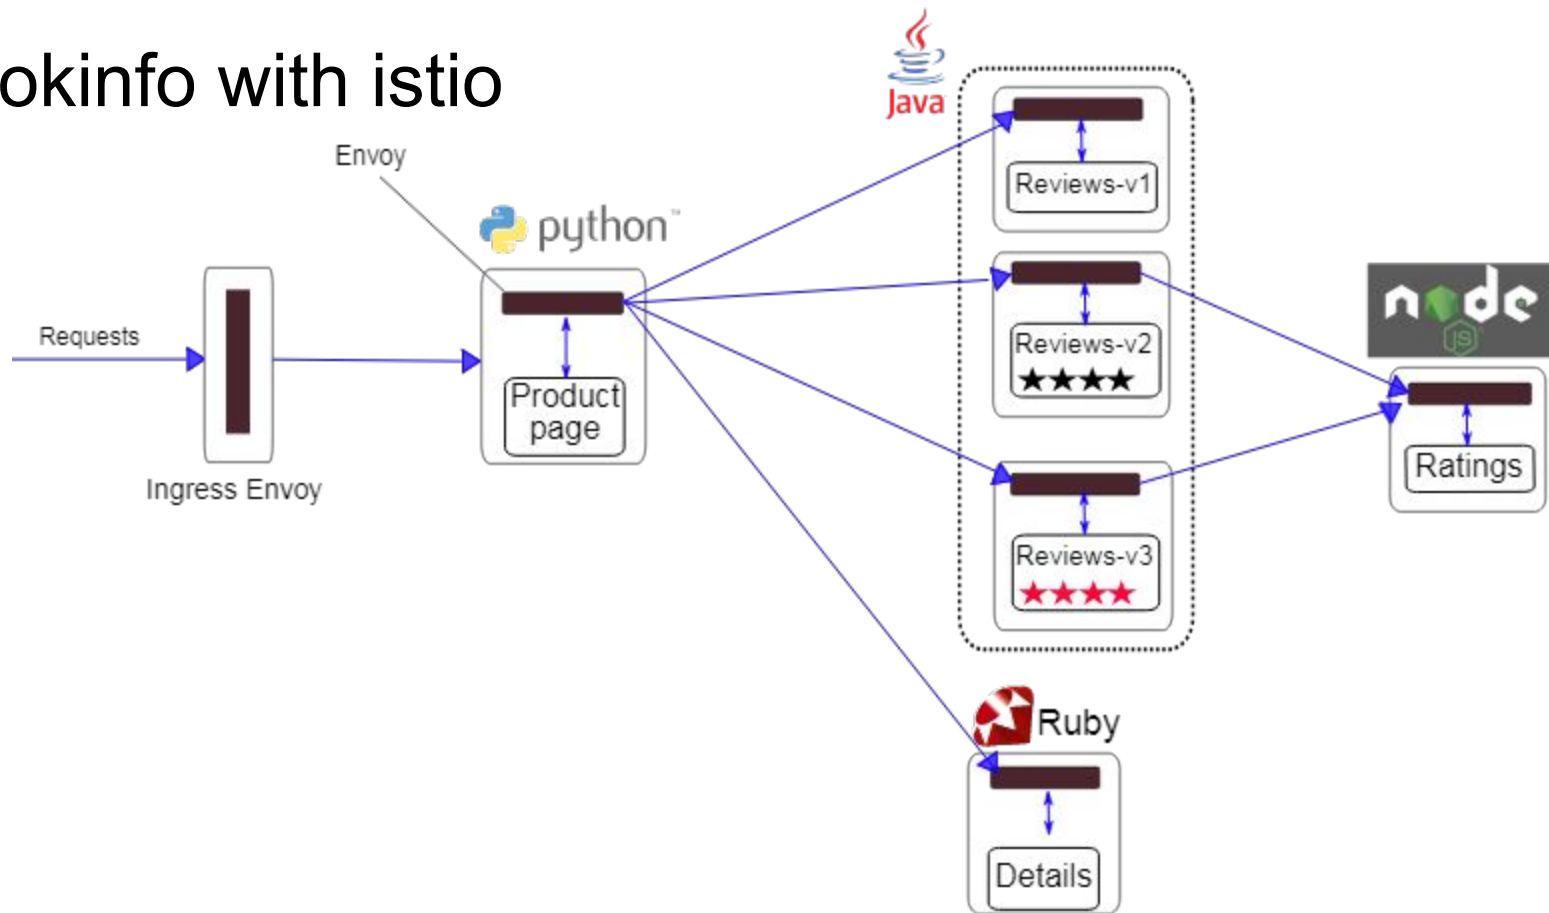
version: v2

Demo usage

Bookinfo



Bookinfo with istio



Splitting traffic between versions

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

name: reviews

spec:

hosts:

- reviews

http:

- route:

- destination:

- host: reviews

- subset: v1

- weight: 75

- destination:

- host: reviews

- subset: v2

- weight: 25

Timeouts

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: ratings

spec:

 hosts:

 - ratings

 http:

 - route:

 - destination:

 host: ratings

 subset: v1

 timeout: 10s

Timeouts and retries

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: ratings

spec:

 hosts:

 - ratings

 http:

 - route:

 - destination:

 host: ratings

 subset: v1

 retries:

 attempts: 3

 perTryTimeout: 2s

Injecting faults (10% and 5 second delay)

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: ratings

spec:

 hosts:

 - ratings

 http:

 - fault:

 delay:

 percent: 10

 fixedDelay: 5s

 route:

 - destination:

 host: ratings

 subset: v1

Injecting faults (10% and http 400 code)

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: ratings

spec:

 hosts:

 - ratings

 http:

 - fault:

 abort:

 percent: 10

 httpStatus: 400

 route:

 - destination:

 host: ratings

 subset: v1

Injecting faults (review-v2 -> rating-v1 delay & abort)

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

name: ratings

spec:

hosts:

- ratings

http:

- match:

- sourceLabels:

- app: reviews

- version: v2

fault:

delay:

fixedDelay: 5s

abort:

percent: 10

httpStatus: 400

route:

- destination:

- host: ratings

- subset: v1

Policies based on subsets

apiVersion: networking.istio.io/v1alpha3

kind: DestinationRule

metadata:

name: reviews

spec:

host: reviews

trafficPolicy:

loadBalancer:

simple: RANDOM

subsets:

- name: v1

labels:

version: v1

- name: v2

labels:

version: v2

trafficPolicy:

loadBalancer:

simple: ROUND_ROBIN

- name: v3

labels:

version: v3

Circuit breakers

apiVersion: networking.istio.io/v1alpha3

kind: DestinationRule

metadata:

 name: reviews

spec:

 host: reviews

 subsets:

 - name: v1

 labels:

 version: v1

 trafficPolicy:

 connectionPool:

 tcp:

 maxConnections: 100

Service entries

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

name: bar-foo-ext-svc

spec:

hosts:

- bar.foo.com

http:

- route:

- destination:

- host: bar.foo.com

- timeout: 10s

Gateway (ingress)

apiVersion: networking.istio.io/v1alpha3

kind: Gateway

metadata:

name: bookinfo-gateway

spec:

selector:

istio: ingressgateway # use istio default controller

servers:

- port:

number: 80

name: http

protocol: HTTP

hosts:

- "*"

Gateway (bind virtual service)

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

 name: bookinfo

spec:

 hosts:

 - "*"

 gateways:

 - bookinfo-gateway

 http:

 - match:

 - uri:

 exact: /productpage

 - uri:

 prefix: /api/v1/products

 route:

 - destination:

 host: productpage

 port:

 number: 9080

开启方式

手工注入sidecar (proxy)

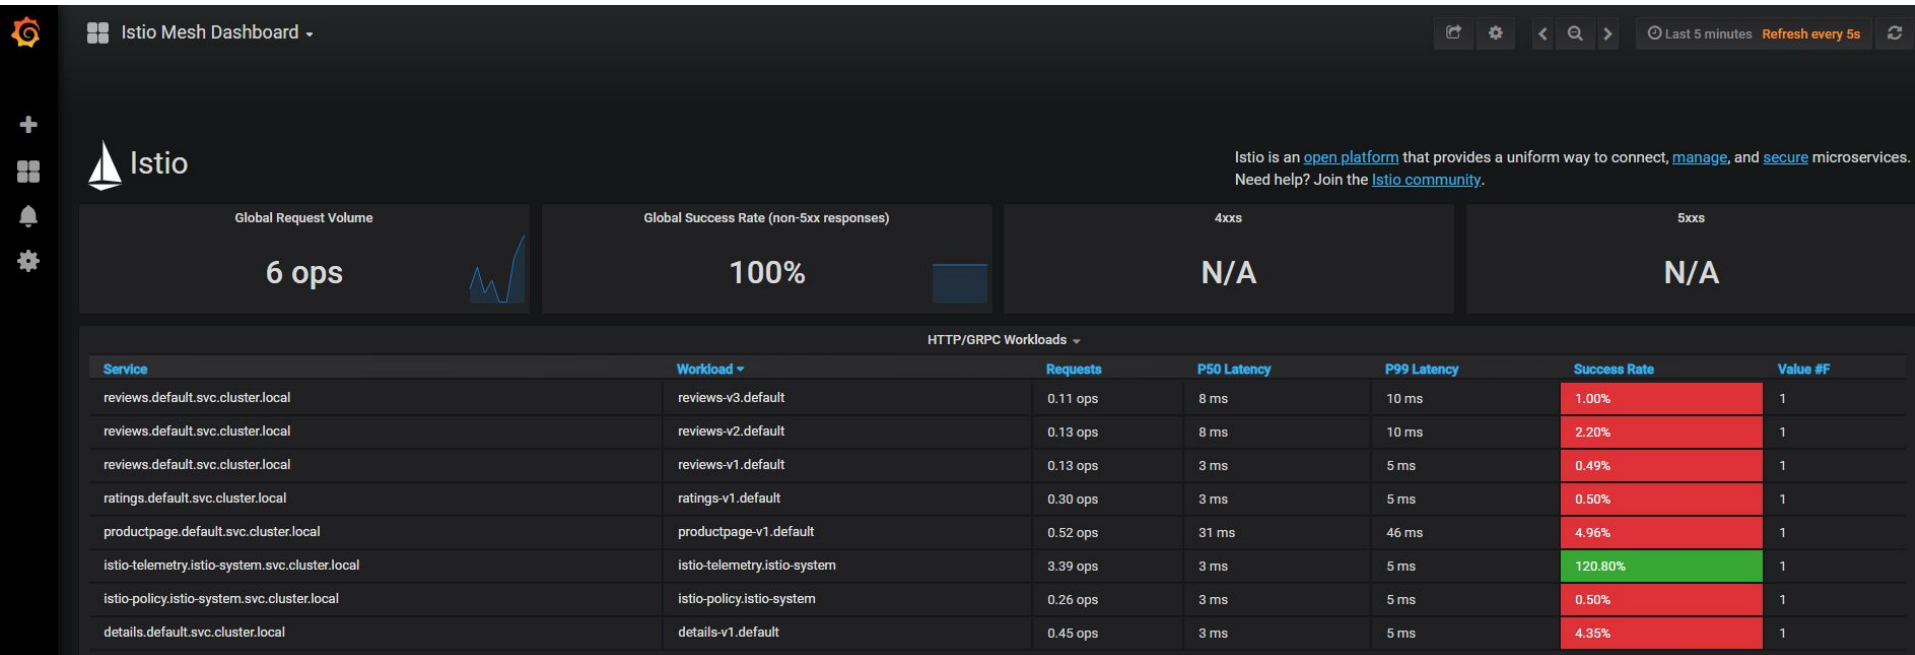
```
istioctl kube-inject -f samples/sleep/sleep.yaml | kubectl apply -f -
```

自动注入sidecar (proxy)

```
kubectl label namespace default istio-injection=enabled
```

```
kubectl apply -f samples/sleep/sleep.yaml
```

Istio grafana





Linkerd

Linkerd介绍

为K8s量身定制的轻量级服务网格方案




Actionable service metrics

Best-in-class observability allows you to monitor *golden metrics*—success rate, request volume, and latency—for every service.




Deep runtime diagnostics

Get a comprehensive suite of diagnostic tools, including automatic service dependency maps and live traffic samples.



Installs in seconds with zero config

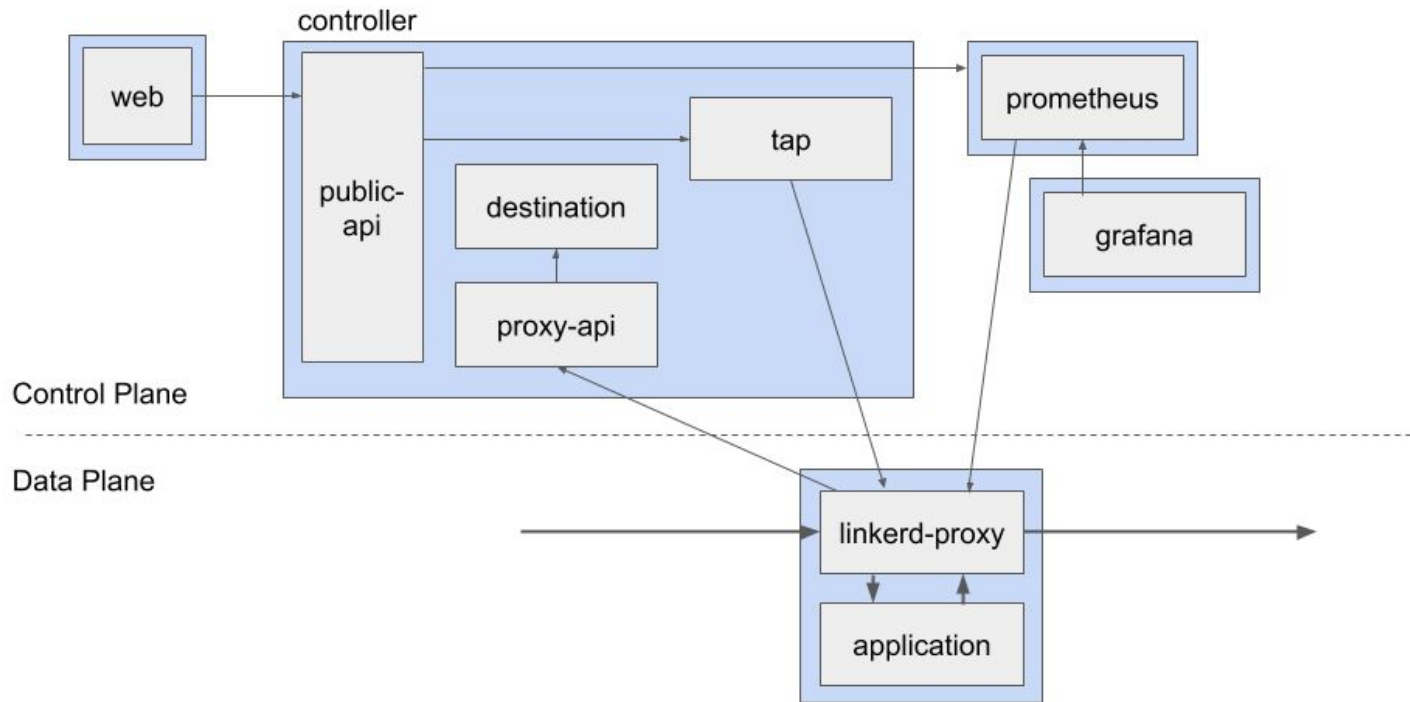
Linkerd's control plane installs into a single namespace, and services can be safely added to the mesh, one at a time.



Simple, minimalist design


No complex APIs or configuration. For most applications, Linkerd will "just work" out of the box.

Linkerd arch



Linkerd UI

Namespace > emojivoto



Deployments

DEPLOYMENT	MESHED	SR	RPS	P50	P95	P99	TLS	DASH
emoji	1/1	100.00%	2	1 ms	1 ms	1 ms	0.0%	
vote-bot	1/1	---	---	---	---	---	---	
voting	1/1	81.67%	1	1 ms	1 ms	1 ms	0.0%	
web	1/1	90.83%	2	2 ms	3 ms	4 ms	0.0%	

Linkerd UI (cont.)

Tap

emojivoto

deployment/web

Start

Hide filters ^

emojivoto

deployment/emoji

Authority

/emojivoto.v1.EmojiService/ListAll

Display requests with this authority

Scheme

Max RPS

Display requests with this scheme

Maximum requests per second to tap (default 100.0)

Display requests with this HTTP method

Current Tap query:

```
linkerd tap deployment/web --namespace emojivoto --to deployment/emoji --to-namespace emojivoto --path /emojivoto.v1.EmojiService/ListAll
```

Direction	Name	Method	Path	Latency	HTTP status	GRPC status	
-	TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/ListAll	1 ms	200	OK
Request Init							
Authority		Path		Scheme	Method	TLS	
emoji-svc.emojivoto:8080		/emojivoto.v1.EmojiService/ListAll		HTTP	POST	no_identity	
Response Init							
HTTP Status		Latency					
200		1 ms					
Response End							
GRPC Status		Latency		Response Length (B)			
OK		41 μs		2,161			
+	TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/ListAll	2 ms	200	OK
+	TO	deploy/emoji	POST	/emojivoto.v1.EmojiService/ListAll	2 ms	200	OK

Linkerd与Istio的对比

Linkerd的也是通过sidecar来中转流量，以提供相应的附件功能，两者都是对业务无侵入的方式实现

Linkerd在监控方面做的更强大点，可观察具体的访问请求

Linkerd在其它方面目前尚未提供太多功能(如流量管理方面)

总结

Istio在服务管理上提供了更多功能(如流量新旧版本分离, A/B测试等)

Linkerd在监控方面做的更强大(具体到查看每个访问的状态)。

通常两者对同一个应用不能同时使用, 由于Istio的复杂性, 及特殊业务场景的应用要求, 初期可以考虑使Istio对某一应用选择性的支持, Linkerd的应用也须作针对性的选择(如对监控要求较高的)。

K8s部署service-mesh相关信息

<http://issue.qianbao-inc.com/SRE/k8s/src/branch/master/service-mesh/istio>

<http://issue.qianbao-inc.com/SRE/k8s/src/branch/master/service-mesh/linkerd2>

Reference

<https://istio.io/>

<https://linkerd.io>

<https://opensource.com/article/18/9/what-istio>

Thank you !